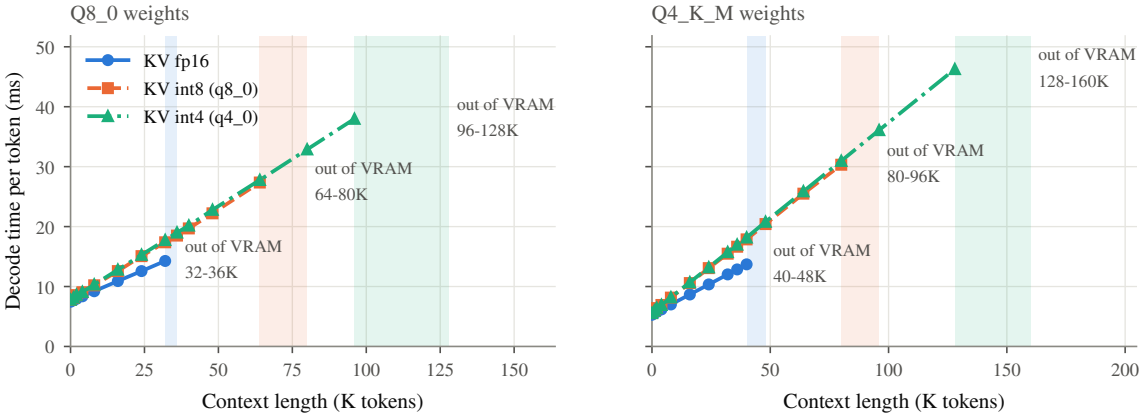
Q8_0 weights
KV fp16
KV int8 (q8_0)
KV int4 (q4_0)
out of VRAM
32-36K
out of VRAM
64-80K
out of VRAM
96-128K
Decode time per token (ms)
Context length (K tokens)
Q4_K_M weights
out of VRAM
40-48K
out of VRAM
80-96K
out of VRAM
128-160K
Context length (K tokens)