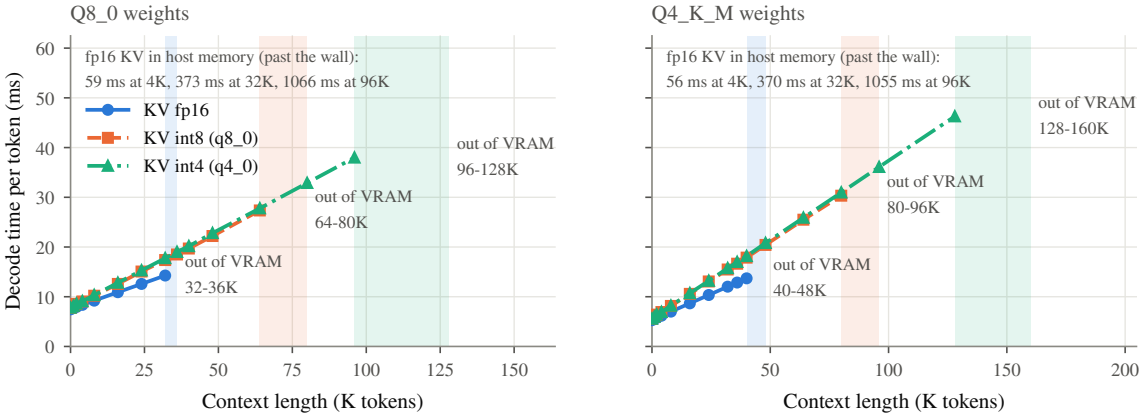
Q8_0 weights
fp16 KV in host memory (past the wall):
59 ms at 4K, 373 ms at 32K, 1066 ms at 96K
KV fp16
KV int8 (q8_0)
KV int4 (q4_0)
out of VRAM
32-36K
out of VRAM
64-80K
out of VRAM
96-128K
Decode time per token (ms)
Context length (K tokens)
Q4_K_M weights
fp16 KV in host memory (past the wall):
56 ms at 4K, 370 ms at 32K, 1055 ms at 96K
out of VRAM
40-48K
out of VRAM
80-96K
out of VRAM
128-160K
Context length (K tokens)